

Library versions with libtool and ABI compatibility

Tuomo Turunen
2019-11-07

What is API? What is ABI?

- API – Application Programming Interface
 - In the scope of this slide set we consider API as **public C and C++ headers** provided by a SW component (i.e. headers available in devel package)
- ABI – Application Binary Interface
 - In the scope of this slide set we consider ABI as the binary interface between a shared library and an application

What is API and ABI compatibility

- A change to API is *backward compatible*, if source code of applications using the API do not have to be changed
 - The rest of this slide set will not talk about API compatibility, but concentrates only on ABI compatibility
- A change to ABI is *backward compatible*, if binaries using the ABI do not have to be recompiled
 - The rest of this slide set will use these terms:
 - ABI compatible: binaries do not have to be recompiled
 - ABI **incompatible**: binaries must be recompiled

Shared library versioning with libtool

- Most RCP subsystems use libtool for building shared libraries
- Libtool adds a version number at the end of a library name based on three values: **current**, **revision** and **age**
- The rules for updating **current**, **revision** and **age** are explained in <https://autotools.io/libtool/version.html>. Copy-pasted also here:
 - Always increase the **revision** value
 - Increase the **current** value whenever an interface has been added, removed or changed
 - Increase the **age** value only if the changes made to the ABI are backward compatible
 - Note #1: if there are no changes made to the ABI (i.e. **current** wasn't increased), then **age** must not be touched
 - Note #2: **revision**, **current** or **age** are never decreased or reset to zero
- There are also other (more complex) guidelines for updating **current**, **revision** and **age**, but in RCP we use this one
- **Warning:** A common mistake is to assume that **current**, **revision** and **age** values map directly into the three numbers at the end of the library name. This is not the case, and indeed, **current**, **revision** and **age** are applied differently depending on the operating system that one is using.
 - Linux: `libxxx.so.{current - age}.{age}.{revision}`
 - FreeBSD: `libxxx.so.{current}`
 - OpenBSD: `libxxx.so.{current}.{revision}`
 - Android: `libxxx.so`

Examples

- Example 1: No interface changes
 - Increase **revision**
 - Do not touch **current** or **age**
- Example 2: ABI compatible change
 - Increase **current**, **revision** and **age**
- Example 3: ABI **incompatible** change
 - Increase **current** and **revision**
 - Do not touch **age**

Shared libraries in Linux

- In Linux a shared library consists of two symbolic links and the actual library:
`libxxx.so -> libxxx.so.{current - age}.{age}.{revision}`
`libxxx.so.{current - age} -> libxxx.so.{current - age}.{age}.{revision}`
`libxxx.so.{current - age}.{age}.{revision}`
- The symbolic link without version number (i.e. `libxxx.so`) is used in development environment for linking applications using the shared library
 - This file is needed only in the development environment and thus belongs to the devel package
- The symbolic link with one version number and the actual library are used by loader when executing binaries
 - These files are needed in both development environment and runtime environment
 - Typically these files belong to `libs` package, but can be also in the base package if there is no need for separate `libs` package
- “*Library soname*” or “*SONAME*” is `libxxx.so.{current - age}`
 - soname is used for both backward-compatibility information and dependency information

Example

- libgenapi: current = 11, revision = 26, age = 7

```
[tuomo@linux:/usr/lib64]
$ ls -l libgenapi.so*
lrwxrwxrwx 1 root root      19 Oct 23 10:35 libgenapi.so -> libgenapi.so.4.7.26
lrwxrwxrwx 1 root root      19 Oct 23 10:35 libgenapi.so.4 -> libgenapi.so.4.7.26
-rwxr-xr-x 1 root root 463160 Oct 23 10:35 libgenapi.so.4.7.26
[tuomo@linux:/usr/lib64]
$ readelf -a libgenapi.so.4.7.26 | grep soname
0x000000000000000e (SONAME)          Library soname: [libgenapi.so.4]
[tuomo@linux:/usr/lib64]
$ rpm -qf libgenapi.so
genapi-devel-1.4.0-1.wf30.x86_64
[tuomo@linux:/usr/lib64]
$ rpm -qf libgenapi.so.4
genapi-1.4.0-1.wf30.x86_64
[tuomo@linux:/usr/lib64]
$ rpm -qf libgenapi.so.4.7.26
genapi-1.4.0-1.wf30.x86_64
```

Why library versions are important?

- There are multiple reasons to have correctly versioned libraries
- In RCP the most important reason is RPM dependency tracking
 - When a binary rpm is built, all shared libraries used by the rpm and provided by the rpm are (automatically) listed as dependencies
 - For example:
`shareddatalayer-libs-3.7.8-1.wf30.x86_64.rpm Provides libSDL.so.8()(64bit)`
`trafficmanager-1.0.4-4.wf30.x86_64.rpm Requires libSDL.so.8()(64bit)`
 - The above dependency means that installing trafficmanager 1.0.4-4 requires also compatible shareddatalayer-libs
- If shared library ABI is changed in incompatible manner, the library version changes and RPM notices the conflict – to fix the conflict, rpms depending on the old library version must be rebuilt
- If library versions are missing or are updated incorrectly, RCP build will not work!
 - This wasn't a problem with the legacy RCP build, because everything was always rebuilt
 - This is a major problem with Koji build, where rpms are rebuilt only when needed

Example

- genapi 1.2.2-1 provides libgenapi.so.3 (libgenapi.so.3.6.24)
- minstarter 0.1.9-1 depends on libgenapi.so.3
- genapi 1.3.0-1 has interface change, which is ABI **incompatible** (but API compatible)
 - Library name is changed to libgenapi.so.4 (libgenapi.so.4.6.25)
- minstarter has to be rebuilt
 - minstarter 0.1.9-2 depends on libgenapi.so.4
- genapi 1.4.0-1 has interface change, which was ABI compatible
 - Library name doesn't change; it is still libgenapi.so.4 (libgenapi.so.4.7.26)
- There is no need to rebuild minstarter; it is still minstarter 0.1.9-2

How to determine if interface change is ABI compatible?

- Sometimes it is not easy to determine whether an interface change is ABI compatible
 - C++ standard doesn't say much about binary interface
 - ABI compatibility cannot be proven by testing (but ABI **incompatibility** can)
- ABI documents are long and complex:
 - <http://static.coldattic.info/restricted/science/syrcose09/cppbincomp.pdf>
 - <http://itanium-cxx-abi.github.io/cxx-abi/abi.html>
- Rules of thumb:
 - If public headers are not touched, then the change is ABI compatible
 - Adding new interfaces is ABI compatible
 - Removing existing interfaces is ABI **incompatible**
 - Changing existing interfaces depends on the change – see the examples in the following slides

Adding new interfaces

- Adding new macros, constants, structs, classes, enums and functions is ABI compatible

```
+ #define NEW_MACRO ... ✓  
+ constexpr int NEW_CONSTANT = ...; ✓  
+ struct NewStruct { ... }; ✓  
+ class NewClass { ... }; ✓  
+ enum class NewEnum { ... }; ✓  
+ void newFunction(int a); ✓
```

Removing existing interfaces

- Removing existing macros, constants, structs, classes, enums and functions is ABI **incompatible**

```
- #define EXISTING_MACRO ...  
- constexpr EXISTING_CONSTANT = ...;  
- struct ExistingStruct { ... };  
- class ExistingClass { ... };  
- enum class ExistingEnum { ... };  
- void existingFunction(int a);
```

Changing existing macros

- Changing existing macros is ABI **incompatible**, because macros are evaluated at compile time

```
- #define MY_BIT_MASK 0x7  
+ #define MY_BIT_MASK 0xF
```



Changing existing constants

- Changing existing constants is ABI **incompatible**, because constexpr statements are evaluated at compile time

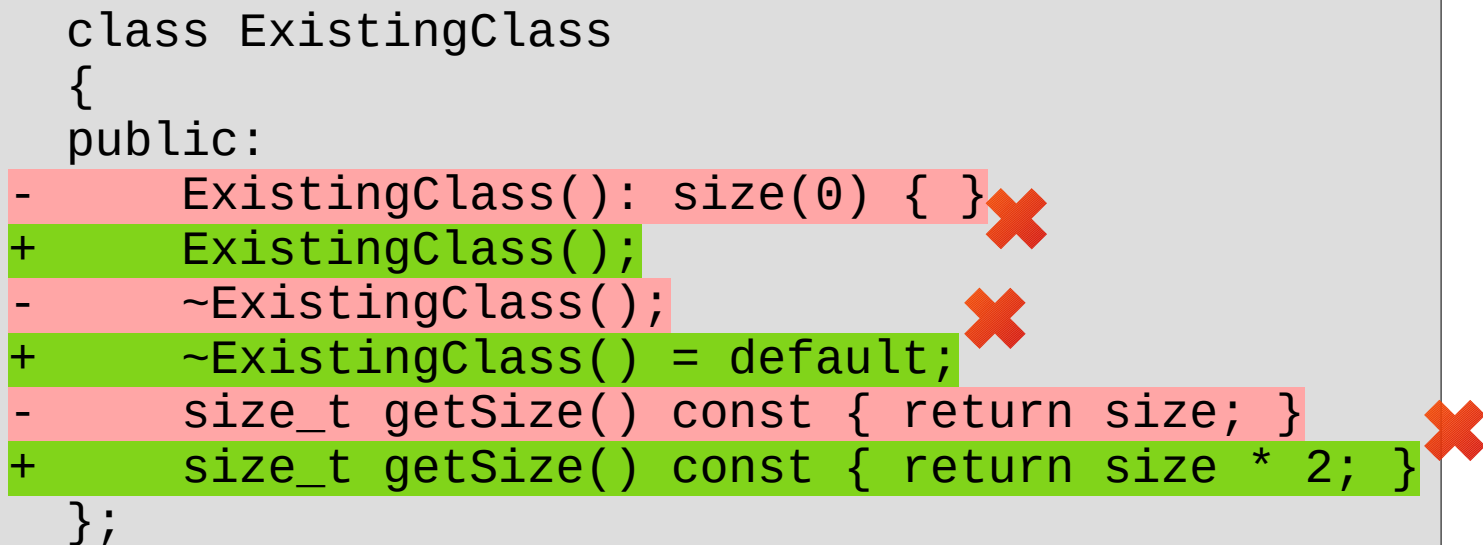
```
- constexpr size_t MAX_SIZE = 32;  
+ constexpr size_t MAX_SIZE = 64;
```



Changing existing inline functions

- Changing existing inline functions, changing existing non-inline functions to inline or changing existing inline functions to non-inline is ABI **incompatible**, because inline functions are evaluated at compile time
- This applies also to the default constructors, destructors, copy constructors, copy operators, move constructors and move operators either implicitly or explicitly created by compiler

```
class ExistingClass
{
public:
-   ExistingClass(): size(0) { }
+   ExistingClass();
-   ~ExistingClass();
+   ~ExistingClass() = default;
-   size_t getSize() const { return size; }
+   size_t getSize() const { return size * 2; }
};
```

The diagram shows a C++ class definition for ExistingClass. The original code (highlighted in pink) is being replaced by new code (highlighted in green). Red 'X' marks are placed at the end of the lines being removed or added to indicate that these changes are ABI incompatible. The changes include: removing the constructor ExistingClass(): size(0) { }, adding a new constructor ExistingClass();, removing the destructor ~ExistingClass();, adding a new destructor ~ExistingClass() = default;, removing the const method getSize() returning size, and adding a new const method getSize() returning size * 2.

Changing existing functions

- Changing existing function return value type is ABI **incompatible**
- Changing existing function parameter type is ABI **incompatible**
- Adding parameter with default value to existing function is ABI **incompatible**

```
- uint32_t getBits();  
+ uint64_t getBits();  
- void print(Object object);  
+ void print(const Object& object);  
- Configuration getConfiguration();  
+ Configuration getConfiguration(bool validate = false);
```


Changing existing member functions

- Adding or removing `const` keyword is ABI **incompatible**
- Adding or removing `virtual` keyword is ABI **incompatible**
- Adding or removing `noexcept` keyword is ABI **incompatible**

```
class ExistingClass
{
public:
-   void draw(Canvas& canvas) const;
+   virtual void draw(Canvas& canvas) const;
-   size_t getSize();
+   size_t getSize() const;
-   void clear();
+   void clear() noexcept;
};
```



Adding or removing member variables

- Adding or removing (private, protected or public) member variable from a struct or a class is ABI **incompatible**, because this changes the size of the struct or class

```
class ExistingClass
{
    ...
private:
    int x;
+   int y; ❌
};
```

Reordering public member variables

- Reordering public or protected member variables in a struct or a class is ABI **incompatible**, because this changes the variable offset in the struct of class
- This applies also to private member variables used by inline functions

```
struct ExistingStruct
{
-   int y; ✖
    int x;
+   int y;
};
```

Adding new non-virtual member functions


- Adding new non-virtual (private, protected or public) member function is ABI compatible
- The order of non-virtual functions is irrelevant

```
class ExistingClass
{
public:
    void oldFunction();
+   void newFunction(); ✓
};
```


Adding new virtual member functions

- Adding new virtual (private, protected or public) member function is ABI compatible, if the function is added as the last virtual function in the class
- Adding new virtual member function between or before existing virtual functions is ABI **incompatible**
- The order of virtual functions is very important!

```
class ExistingClass
{
public:
    virtual void oldFunction();
+   virtual void newFunction();
};
```



```
class ExistingClass
{
public:
+   virtual void newFunction();
    virtual void oldFunction();
};
```



Adding new static member variables or functions

- Adding new static (private, protected or public) member variable or function is ABI compatible

```
class ExistingClass
{
+   static int variable; ✓
+   static void function(); ✓
};
```

Changing class hierarchy

- In general, changing class hierarchy in anyway is ABI **incompatible**

```
- class ExistingClass
+ class ExistingClass: public NewBaseClass
{
    ...;
};
```



Adding new enum values

- Adding a new enum value without changing any of the existing enum values is ABI compatible (as long as the new value fits into the underlying type, which is by default `int`)
- Changing an existing enum value is ABI **incompatible**

```
enum class ExistingEnum
{
    RED,
    GREEN,
+   BLUE, ✓
};
```

```
enum class ExistingEnum
{
    RED,
+   BLUE, ✗
    GREEN,
};
```


Techniques to avoid incompatible ABI changes

- API (i.e. the public headers) must be designed so that it reveals as little as possible about the actual implementation – this must be the design from the very beginning
- Avoid macros in public headers
- Avoid inline functions in public headers
 - Don't inline functions for optimization, unless you have measured performance and measurement proves that inlining functions is the solution
- Always add new enums at the end of the list
- Avoid functions with default value parameters in public headers
 - Use function overloading instead
- Do not expose class implementation in public headers
 - Forward declaration pattern
 - PImpl pattern
 - Interface pattern with abstract classes

Forward declaration pattern

- Forward declaration pattern is typically used in C APIs
- The idea is to expose only struct name in public header, but nothing about its implementation – application can use the struct only via the functions declared in the public header

public_header.h

```
struct my_struct; // forward declaration

struct my_struct *create_my_struct(...);
void use_my_struct(struct my_struct *);
void destroy_my_struct(struct my_struct *);
```

implementation.c

```
struct my_struct
{
    int x;
    int y;
};
```

PImpl pattern

- *Private Implementation* pattern is basically C++ version of the forward declaration – the real implementation is not in the public header

public_header.hpp

```
class MyClass
{
public:
    MyClass();
    ~MyClass();
    void function();
private:
    class Impl; // forward declaration
    std::unique_ptr<Impl> impl;
};
```

implementation.cpp

```
class MyClass::Impl
{
public:
    void function();
private:
    ...
};

MyClass::MyClass(): impl(new Impl()) { }

MyClass::~MyClass() { }

void MyClass::function() { impl->function(); }
```

Interface pattern

- Exposing only abstract classes in public headers not only completely hides the implementation from application but also allows *decoupling* as defined by Dependency Inversion Principle (DIP)

public_header.hpp

```
class MyClass
{
public:
    MyClass() = default;
    virtual ~MyClass() = default;
    virtual void function() = 0;
    static std::unique_ptr<MyClass> create();
};
```

implementation.cpp

```
class MyClassImpl: public MyClass
{
public:
    void function() override;
private:
    ...
};

std::unique_ptr<MyClass> MyClass::create()
{
    return std::make_unique(MyClassImpl);
}
```

Incompatible ABI changes via FCI

- Doing ABI **incompatible** changes via FCI is explained in <https://confluence.int.net.nokia.com/display/RCP/FCI+support+for+koji>
- Briefly:
 - If a subsystem has ABI **incompatible** change, then add the source code repository name to "Incompatible ABI Change" box
 - Determining whether ABI has no changes, has ABI compatible changes or has ABI **incompatible** changes is up to the **developer** – this is not done automatically by FCI
 - Updating library version numbers is up to the **developer** – this is not done automatically by FCI

Links

- Autotools Mythbuster – Library Versioning:
<https://autotools.io/libtool/version.html>
- soname
<https://en.wikipedia.org/wiki/Soname>
- Binary Compatibility of Shared Libraries Implemented in C++ on GNU/Linux Systems
<http://static.coldattic.info/restricted/science/syrcoese09/cppbincomp.pdf>
- Itanium C++ ABI
<http://itanium-cxx-abi.github.io/cxx-abi/abi.html>
- FCI support for koji
<https://confluence.int.net.nokia.com/display/RCP/FCI+support+for+koji>